

DataFrame Basic Functions and Groupby

Johnny Zhang

Changing Column Names

- For either development or reporting purposes you will often want to rename columns.

```
# Import data into a DataFrame.  
  
path = "/Users/pm/Desktop/DayDocs/2019_2020/PythonForDataAnalytics/workingData/babysamp-98.txt"  
  
df = pd.read_csv(path, skiprows=1, sep='\t', names=('MomAge', 'DadAge', 'MomEduc', 'MomMarital', 'numlive',  
"dobmm", 'gestation', 'sex', 'weight', 'prenatalstart', 'orig.id', 'preemie'))  
  
# Rename the columns so they are more reader-friendly.  
  
df = df.rename({'MomAge': 'Mom Age', 'DadAge':'Dad Age',  
               'MomEduc':'Mom Edu', 'weight':'Weight'}, axis=1) # new method  
  
# Show all columns.  
  
pd.set_option('display.max_columns', None)  
  
# Increase number of columns that display on one line.  
  
pd.set_option('display.width', 1000)  
  
print(df.head())
```

Exercise 1

Analyzing Frequencies (value_counts)

- With the value_counts function you can get a quick understanding of the column ranges and frequency for each value.
- count vs value_counts
 - value_counts: Category & count for each category
 - It is specifically for Series (single columns), but it can be used on a DataFrame by calling it on a specific column (`(df['column_name'].value_counts())`).
 - count: total number of rows (It does not count NaN or None Value)

`Series.value_counts(normalize=False, sort=True, ascending=False,
bins=None, dropna=True)`

value_counts() vs. count()

```
# Sample Series  
data = pd.Series([1, 2, 2, 3, 3, 3, 4])
```

```
# Using value_counts on Series  
print(data.value_counts())
```

```
3    3  
2    2  
1    1  
4    1  
dtype: int64
```

```
data = pd.Series([1, 2, 2, 3, None, 3, 3, 4, None])  
print(data.value_counts(dropna=False))
```

```
3.0    3  
2.0    2  
1.0    1  
4.0    1  
NaN    2  
dtype: int64
```

It ignores NaN values by default but can include them if specified using the `dropna=False` parameter.

```
import pandas as pd
# Sample DataFrame
df = pd.DataFrame({
    'A': [1, 2, 2, 3, 4, 4, 4, 5],
    'B': ['a', 'b', 'b', 'c', 'd', 'd', 'e', 'f']
})
# Get unique values from column 'A'
unique_values_A = df['A'].unique()
print(unique_values_A)

# Get unique values from column 'B'
unique_values_B = df['B'].unique()
print(unique_values_B)

print(list(set(df["B"])))
```

```
[1 2 3 4 5]
['a' 'b' 'c' 'd' 'e' 'f']
['d', 'f', 'b', 'a', 'c', 'e']
```

```

print("\nTOP FREQUENCY FIRST")
print(df['Mom Age'].value_counts())

print("\nLOWEST FREQUENCY FIRST")
print(df['Mom Age'].value_counts(ascending=True))

print("\nFREQUENCY SORTED by MOTHER AGE")
print(df['Mom Age'].value_counts().sort_index())

```

	TOP FREQUENCY FIRST	LOWEST FREQUENCY FIRST	FREQUENCY SORTED by MOTHER AGE
20	16	14	14 1
22	13	39	15 3
24	13	42	16 2
27	12	38	17 3
23	12	16	18 6
21	11	17	19 10
19	10	40	20 16
29	9	15	21 11
33	9	41	22 13
30	9	35	23 12
26	9	37	24 13
36	8	34	25 8
25	8	32	26 9
31	7	28	27 12
18	6	18	28 6
28	6	31	29 9
32	6	25	30 9
34	6	36	31 7
37	6	26	32 6
35	5	30	33 9
41	3	33	34 6
15	3	29	35 5
40	3	19	36 8
17	3	21	37 6
16	2	23	38 2
38	2	27	39 1
42	1	24	40 3
39	1	22	41 3
14	1	20	42 1

Exercise 2

The `sort_index()` method in pandas is used to sort the rows or columns of a DataFrame or Series based on their index labels. You can sort the data in ascending or descending order.

```
import pandas as pd

# Sample DataFrame
df = pd.DataFrame({
    'A': [10, 20, 30, 40],
    'B': [50, 60, 70, 80]
}, index=['d', 'a', 'c', 'b'])

print("Original DataFrame:")
print(df)

# Sorting by index (rows) in ascending order
df_sorted = df.sort_index()
print("\nSorted by index (ascending):")
print(df_sorted)
```

Original DataFrame:

	A	B
d	10	50
a	20	60
c	30	70
b	40	80

Sorted by index (ascending):

	A	B
a	20	60
b	40	80
c	30	70
d	10	50

```
df_sorted_columns = df.sort_index(axis=1)
print("\nSorted by columns index (ascending):")
print(df_sorted_columns)
```

Sorted by columns index (ascending):

	A	B
a	20	60
b	40	80
c	30	70
d	10	50

axis :

- `axis=0` (default) sorts the rows by index labels.
- `axis=1` sorts the columns by column labels.

ascending :

- Set to `True` (default) for ascending order.
- Set to `False` for descending order.

Sorting a DataFrame (sort_values)

- The sort_values method allows us to sort DataFrame rows by multiple columns values in either ascending or descending order;

```
dfSorted = df.sort_values(['Mom Age', 'Weight'], ascending=[False, True])
```

```
print(dfSorted)
```

	Mom	Age	Dad	Age	Mom	Edu	MomMarital	numLive	dobmm	gestation	sex	Weight	prenatalstart	orig.id	preemie
100		42		43.0		15.0		1	4	4	M	4224	2.0	2582947	False
39		41		47.0		12.0		1	3	7	M	2182	2.0	1622541	False
8		41		39.0		14.0		1	0	11	M	2834	2.0	2481971	True
130		41		41.0		16.0		1	2	7	F	4000	1.0	2267907	False
115		40		32.0		17.0		1	1	10	M	2268	1.0	51593	True
..
56		16		NaN		8.0		2	0	7	M	1985	2.0	1247735	False
95		15		NaN		8.0		2	0	2	F	2552	5.0	3145461	True
90		15		NaN		8.0		2	0	7	F	2892	1.0	945363	False
73		15		20.0		NaN		2	0	5	F	3459	5.0	1316782	False
160		14		NaN		8.0		2	0	10	F	2977	6.0	1123508	False

```
import pandas as pd
import matplotlib.pyplot as plt

# Import data into a DataFrame.
path = "/your/file/full/path/filename.file_extension"
df = pd.read_csv(path, skiprows=1,
                  sep='\t',
                  names=('MomAge', 'DadAge', 'MomEduc', 'MomMarital', 'numlive',
                         "dobmm", 'gestation', 'sex', 'weight', 'prenatalstart',
                         'orig.id', 'preemie'))

# Rename the columns so they are more reader-friendly.
df = df.rename({'MomAge': 'Mom Age', 'DadAge':'Dad Age',
                'MomEduc':'Mom Edu', 'weight':'Weight'}, axis=1) # new method
# Show all columns.
pd.set_option('display.max_columns', None)

# Increase number of columns that display on one line.
pd.set_option('display.width', 1000)

# Sort by descending mother age and then by ascending weight.
dfSorted = df.sort_values(['Mom Age', 'Weight'], ascending=[False, True])
print(dfSorted)
```

Exercise 3

Filtering a DataFrame

It is possible to query a DataFrame with **conditional expressions** and compound conditions:

```
# Compound conditions require single '&' for 'AND'  
# and single '|' for 'OR.  
result = df[(df['DadAge']>40) & (df['MomAge'] > 30)]  
result = df[(df['DadAge']>40) | (df['MomAge'] > 30)]  
  
import matplotlib.pyplot as plt  
import pandas as pd  
# Import data into a DataFrame.  
path = "/Users/pm/Desktop/DayDocs/2019_2020/PythonForDataAnalytics/workingData/babysamp-  
98.txt"  
df = pd.read_csv(path, skiprows=1,  
                 sep='\t',  
                 names=('MomAge', 'DadAge', 'MomEduc', 'MomMarital', 'numlive',  
                        "dobmm", 'gestation', 'sex', 'weight', 'prenatalstart',  
                        'orig.id', 'preemie'))  
# Show all columns.  
pd.set_option('display.max_columns', None)  
# Increase number of columns that display on one line.  
pd.set_option('display.width', 1000)  
# Compound conditions require single '&' for 'AND' and single '|' for 'OR.  
resultDf = df[(df['DadAge']>=40) & (df['MomAge'] >= 40)]  
print(resultDf)
```

	MomAge	DadAge	MomEduc	MomMarital	numlive	dobmm	gestation	sex	weight	prenatalstart	orig.id	preemie
39	41	47.0	12.0	1	3	7	38	M	2182	2.0	1622541	False
100	42	43.0	15.0	1	4	4	40	M	4224	2.0	2582947	False
130	41	41.0	16.0	1	2	7	40	F	4000	1.0	2267907	False
186	40	44.0	12.0	1	1	3	36	M	2693	2.0	2990722	True

Numeric DataFrame Summaries

Function	Description
count()	Number of non-null observations
sum()	Sum of values
mean()	Mean of values
median()	Arithmetic median of values
min()	Minimum
max()	Maximum
std()	Unbiased standard deviation

Syntax:

```
print(df['MomAge'].mean())
```

Exercise 4 (3 marks)

```
Count: 200
Min: 14
Max: 42
Mean: 26.585
Median: 26.0
Standard Deviation: 6.484056596741636
```

count() vs. value_counts()

```
import numpy as np
import pandas as pd
company=['A','B','C','D']
data=pd.DataFrame({'Company':[company[x] for x in np.random.randint(0,len(company),10)],\
                   'Salary':np.random.randint(5,50,10),\
                   'Age':np.random.randint(15,50,10)})
print(data)
print('count function: ',data['Company'].count())
print('value_counts: \n',data['Company'].value_counts())
```

	Company	Salary	Age
0	C	33	44
1	B	19	42
2	D	44	45
3	D	17	20
4	D	41	39
5	D	17	47
6	D	37	37
7	C	30	47
8	B	26	20
9	C	35	17

```
count function: 10
value_counts:
   D    5
   C    3
   B    2
Name: Company, dtype: int64
```

```
import numpy as np

company=list("ABCD")

newArr=[]

for x in np.random.randint(0, len(company), 10):
    newArr.append(company[x])
print(newArr)
```

Grouping on Columns(groupby())

```
import numpy as np
from pandas import DataFrame

Company=['A','B','C']
data=DataFrame({'Company':[Company[index] for index in np.random.randint(0, len(Company), 10)], \
                 'Salary':np.random.randint(5,50,10), \
                 'Age': np.random.randint(15,50,10)})
print(data)
```

	Company	Salary	Age
0	A	30	25
1	A	20	17
2	B	49	44
3	A	48	17
4	A	29	28
5	B	35	24
6	C	6	25
7	C	21	43
8	B	35	33
9	B	34	17

```
import numpy as np
from pandas import DataFrame

Company=['A','B','C']
data=DataFrame({'Company':[Company[index] for index in np.random.randint(0, len(Company), 10)],\
                 'Salary':np.random.randint(5,50,10),\
                 'Age': np.random.randint(15,50,10)})
print(data)
```

	Company	Salary	Age
0	A	36	37
1	B	49	47
2	C	11	42
3	C	18	40
4	C	5	39
5	B	6	19
6	B	39	44
7	A	24	21
8	A	46	28
9	B	39	24

```
print(data.groupby('Company'))
```

```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x000002761E822BE0>
```

```
print(list(data.groupby('Company'))))
```

[('A', DataFrame(Company, Salary, Age)), ('B', DataFrame(Company, Salary, Age)), ('C', DataFrame(Company, Salary, Age))]

The output shows three groups of data, each represented by a DataFrame. The first group, labeled 'A', has 3 rows. The second group, labeled 'B', has 4 rows. The third group, labeled 'C', has 3 rows. Each group contains columns for Company, Salary, and Age.

```
print(type(list(data.groupby('Company'))[0]))
```

```
<class 'tuple'>
```

```
print(list(data.groupby('Company'))[0])
```

```
('A',    Company  Salary  Age
0      A        36     37
7      A        24     21
8      A        46     28)
```

```
print(list(data.groupby('Company'))[1])
```

```
('B',    Company  Salary  Age
1      B        49     47
5      B        6      19
6      B        39     44
9      B        39     24)
```

```
print(list(data.groupby('Company'))[2])
```

```
('C',    Company  Salary  Age
2      C        11     42
3      C        18     40
4      C        5      39)
```

```
| print(list(data.groupby('Company')))
```

```
[('A',    Company  Salary  Age
0      A        36     37
7      A        24     21
8      A        46     28), ('B',    Company  Salary  Age
1      B        49     47
5      B        6      19
6      B        39     44
9      B        39     24), ('C',    Company  Salary  Age
2      C        11     42
3      C        18     40
4      C        5      39)]
```

```
print(list(data.groupby('Company') ['Company']))
```

```
[('A', 0      A
7      A
8      A
Name: Company, dtype: object), ('B', 1      B
5      B
6      B
9      B
Name: Company, dtype: object), ('C', 2      C
3      C
4      C
Name: Company, dtype: object)]
```



```
print(list(data.groupby('Company')[['Company', 'Age']])))
```

```
[('A',    Company  Salary  Age
 0      A        36     37
 1      A        24     21
 2      A        46     28), ('B',    Company  Salary  Age
 3      B        49     47
 4      B        6      19
 5      B        39     44
 6      B        39     24), ('C',    Company  Salary  Age
 7      C        11     42
 8      C        18     40
 9      C        5      39)]
```

```
data
```

	company	salary	age
0	B	26	48
1	B	6	29
2	A	48	32
3	B	5	39
4	C	11	46
5	B	41	27
6	B	12	42
7	B	18	20
8	B	37	19
9	A	32	38

```
list(group)
```

```
[('A',  
  0    30  
  1    20  
  3    48  
  4    29  
  Name: Salary, dtype: int32),  
 ('B',  
  2    49  
  5    35  
  8    35  
  9    34  
  Name: Salary, dtype: int32),  
 ('C',  
  6    6  
  7    21  
  Name: Salary, dtype: int32)]
```

```
group=data.groupby('Company')['Salary'].mean()
```

```
group
```

```
Company  
A    31.75  
B    38.25  
C    13.50  
Name: Salary, dtype: float64
```

Functions:

- count: Number of non-NA values in the group
- sum: Sum of non-NA values
- mean: Mean of non-NA values
- Median: Arithmetic median of non-NA values

Functions:

- Std, var: Unbiased standard deviation and variance
- min, max: Minimum and maximum of non-NA values
- prod: Product of non-NA values
- first, last: First and last non-NA values

Output as Series

```
import pandas as pd

# The data file path and file name need to be configured.
PATH = "/Python/DataSets/"
CSV_DATA = "phone_data.csv"

# Note this has a comma separator.
df = pd.read_csv(PATH + CSV_DATA, skiprows=1, encoding="ISO-8859-1", sep=',',
                  names=('index', 'date', 'duration', 'item', 'month', 'network',
                         'network_type'))

dfStats=df.groupby('network')['index'].count()
print(dfStats)
```

- A Series is essentially a one-dimensional labeled array.
- It can hold data of any type (integers, floats, strings, etc.).
- It is similar to a list or a numpy array, but it comes with labels (called index) for each value.

network	
Meteor	87
Tesco	84
Three	215
Vodafone	215
data	150
landline	42
special	3
voicemail	27
world	7

```
import pandas as pd

# Creating a Series
data = [10, 20, 30, 40]
series = pd.Series(data, index=["a", "b", "c", "d"])

print(series)

a    10
b    20
c    30
d    40
dtype: int64
```

Output as a DataFrame

```
dfStats=df.groupby('network')['index'].count()
```

```
import pandas as pd  
  
# Creating a Series  
data = [10, 20, 30, 40]  
series = pd.Series(data, index=["a", "b", "c", "d"])  
print(series)
```

```
a    10  
b    20  
c    30  
d    40  
dtype: int64
```

```
dfStats = df.groupby('network')['index'].count().reset_index().rename(columns={'index': '# Calls'})
```

network	
Meteor	87
Tesco	84
Three	215
Vodafone	215
data	150
landline	42
special	3
voicemail	27
world	7

	network	index
0	Meteor	87
1	Tesco	84
2	Three	215
3	Vodafone	215
4	data	150
5	landline	42
6	special	3
7	voicemail	27
8	world	7

```
# Creating a DataFrame  
data = {"A": [1, 2, 3], "B": [4, 5, 6]}  
df = pd.DataFrame(data)  
print(df)
```

A	B
1	4
2	5
3	6

	network	# Calls
0	Meteor	87
1	Tesco	84
2	Three	215
3	Vodafone	215
4	data	150
5	landline	42
6	special	3
7	voicemail	27
8	world	7

```
<class 'pandas.core.frame.DataFrame'>
```

```

import pandas as pd

# The data file path and file name need to be configured.
PATH    = "/Users/pm/Desktop/DayDocs/2019_2020/PythonForDataAnalytics/workingData/"
CSV_DATA = "phone_data.csv"

# Note this has a comma separator.
df = pd.read_csv(PATH + CSV_DATA, skiprows=1, encoding = "ISO-8859-1", sep=',',
                  names=('index', 'date', 'duration', 'item', 'month','network',
                         'network_type' ))

# Get count of items per month.
dfStats = df.groupby('network')['index']\n    .count().reset_index().rename(columns={'index': '# Calls'})

# Get duration mean for network groups and convert to DataFrame.
dfDurationMean = df.groupby('network')['duration']\n    .mean().reset_index().rename(columns={'duration': 'Duration Mean'})

# Get duration max for network groups and convert to DataFrame.
dfDurationMax = df.groupby('network')['duration']\n    .max().reset_index().rename(columns={'duration': 'Duration Max'})

# Append duration mean to stats matrix.
dfStats['Duration Mean'] = dfDurationMean['Duration Mean']

# Append duration max to stats matrix.
dfStats['Duration Max'] = dfDurationMax['Duration Max']

```

	network	# Calls	Duration Mean	Duration Max
0	Meteor	87	83.137931	1090.000
1	Tesco	84	164.773810	1234.000
2	Three	215	170.004651	2328.000
3	Vodafone	215	68.697674	1859.000
4	data	150	34.429000	34.429
5	landline	42	438.880952	10528.000
6	special	3	1.000000	1.000
7	voicemail	27	65.740741	174.000
8	world	7	1.000000	1.000

Exercise 5

Exercise 6

Write a program to group phone calls by **network type**. Show network_type, total calls, duration mean, duration max, duration minimum and standard deviation in your output.

	network_type	# Calls	Duration Mean	Duration Max	Duration Min	\
0	data	150	34.429000	34.429	34.429	
1	landline	42	438.880952	10528.000	3.000	
2	mobile	601	120.457571	2328.000	1.000	
3	special	3	1.000000	1.000	1.000	
4	voicemail	27	65.740741	174.000	1.000	
5	world	7	1.000000	1.000	1.000	

	Duration Standard Deviation
0	0.000000
1	1631.415609
2	285.077689
3	0.000000
4	44.294984
5	0.000000

Exercise 7 Using the baby sample, group on either male or female and show the maximum weight, minimum weight, and mean weight. Show your program here.

	sex	Max_Weight	Min_Weight	Mean_Weight
0	F	4825	907	3265.625000
1	M	4593	1671	3299.650485

Generating a Calculated Column

If the same calculation can be applied to each row of a column you can do it all at once.

```
import pandas as pd
```

```
# Create data set.  
dataSet = { 'Fahrenheit': [85,95,91] }
```

```
# Create dataframe with data set and named columns.  
# Column names must match the dataSet properties.  
df = pd.DataFrame(dataSet, columns= ['Fahrenheit'])
```

```
df['Celsius'] = (df['Fahrenheit']-32)*5/9  
# Show DataFrame  
print(df)
```

Exercise 8

```
import pandas as pd  
  
# Create data set.  
dataSet = { 'Fahrenheit': [85,95,91] }  
  
# Create dataframe with data set and named columns.  
# Column names must match the dataSet properties  
df = pd.DataFrame(dataSet, columns= ['Fahrenheit'])  
  
# Show DataFrame  
print(df)
```

	Fahrenheit
0	85
1	95
2	91

```
import pandas as pd  
  
# Create data set.  
dataSet = { 'Fahrenheit': [85,95,91] }  
  
# Create dataframe with data set and named columns.  
# Column names must match the dataSet properties.  
df = pd.DataFrame(dataSet, columns= ['Fahrenheit'])
```

```
df['Celsius'] = (df['Fahrenheit']-32)*5/9  
# Show DataFrame  
print(df)
```

	Fahrenheit	Celsius
0	85	29.444444
1	95	35.000000
2	91	32.777778

Grouping on Multiple Columns

```
import pandas as pd
```

```
# The data file path and file name need to be configured.
```

```
PATH    = "/Users/pm/Desktop/DayDocs/2019_2020/PythonForDataAnalytics/workingData/"  
CSV_DATA = "phone_data.csv"
```

```
# Note this has a comma separator.
```

```
df = pd.read_csv(PATH + CSV_DATA, skiprows=1, encoding = "ISO-8859-1", sep=',',  
                  names=('index', 'date', 'duration', 'item', 'month','network',  
                         'network_type' ))
```

```
df2 = df.groupby(['network','item'])['duration'].mean().reset_index()
```

```
print(df2)
```

↑
List

Exercise 9

	network	item	duration
0	Meteor	call	133.333333
1	Meteor	sms	1.000000
2	Tesco	call	194.760563
3	Tesco	sms	1.000000
4	Three	call	284.875000
5	Three	sms	1.000000
6	Vodafone	call	221.530303
7	Vodafone	sms	1.000000
8	data	data	34.429000
9	landline	call	438.880952
10	special	sms	1.000000
11	voicemail	call	65.740741
12	world	sms	1.000000

Updating on DataFrame Cell at a Time

Data frame with original data in Celsius:

```
      City    Temperature
0    Mumbai          23.0
1   Beijing         -11.0
['City', 'Temperature']
```

Dataframe after changed to Fahrenheit:

```
      City    Temperature
0    Mumbai          73.4
1   Beijing         12.2
```

.iat vs .loc

The line of code needed to update the cell uses the `iat[]` reference:

```
df.iat[i, tempColumnPosition] = celsius*9.0/5.0 + 32
```

```
#Convert Columns to A list
columnList = list(df.keys())
```

```
for i in range(0, len(columnList)):
```

```
    if(columnList[i]==columnName):
```

```
        columnPosition = i
```

```
        break # Exit the loop.
```

```
return columnPosition
```

Exercise 10

Dataframe after changed to Fahrenheit:

	<u>City</u>	<u>Temperature</u>
0	Mumbai	73.4
1	Beijing	12.2

Dynamically adjust Example 4 to change the city values, one cell at a time, so the updated data in the data frame becomes:

	City	Temperature
0	The city of Mumbai	73.4
1	The city of Beijing	12.2